

Regularizing Bayesian linear models with an informative prior on R^2

Ben Goodrich* Jonah Gabry* Andrew Gelman* Matthew Stephens†

25 February 2016

Abstract

In this paper we derive an approach for expressing prior beliefs about the location of the R^2 , the familiar proportion of variance in the outcome variable that is attributable to the predictors under a linear model. In particular, when there are many predictors relative to the number of observations we would expect the joint prior derived here to work better than placing independent, heavy-tailed priors on the coefficients, which is standard practice in applied Bayesian data analysis but neither reflects the beliefs of the researcher nor conveys enough information to stabilize all the computations.

1 Introduction

Fully making Bayesian estimation of linear models routine for applied researchers requires prior distributions that work well for any data generated according to the assumptions of the likelihood function. Most Bayesian approaches require the researcher to specify a joint prior distribution for the regression coefficients (and the intercept and error variance), but most applied researchers have little inclination to specify all these prior distributions thoughtfully and take a short-cut by specifying one prior distribution that is taken to apply to all the regression coefficients as if they were independent of each other (and the intercept and error variance).

In this paper we derive and demonstrate an approach for expressing prior beliefs about the location of the R^2 , the familiar proportion of variance in the outcome variable that is attributable to the predictors under a linear model. Since the R^2 is a well-understood bounded scalar, it is easy to specify prior information about it. In particular, when there are many predictors relative to the number of observations we would expect the joint prior

*Columbia University

†University of Chicago

derived here to work better than placing independent, heavy-tailed priors on the coefficients, which neither reflects the beliefs of the researcher nor conveys enough information to stabilize all the computations.

2 Likelihood

The likelihood for one observation y_i under a linear model can be written as the conditionally normal density

$$f(y_i | \mu_i, \sigma_\epsilon) = \frac{1}{\sigma_\epsilon \sqrt{2\pi}} \exp \left\{ -\frac{1}{2} \left(\frac{y_i - \mu_i}{\sigma_\epsilon} \right)^2 \right\},$$

where $\mu_i = \alpha + \mathbf{x}_i^\top \boldsymbol{\beta}$ is a linear predictor and σ_ϵ is the standard deviation of the error in predicting the outcome. The likelihood of the entire sample is the product of N individual likelihood contributions and it is well known that the likelihood is maximized when the sum-of-squared residuals is minimized. This occurs when

$$\begin{aligned} \hat{\boldsymbol{\beta}} &= (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}, \\ \hat{\alpha} &= \bar{y} - \bar{\mathbf{x}}^\top \hat{\boldsymbol{\beta}}, \\ \hat{\sigma}_\epsilon^2 &= \frac{1}{N} (\mathbf{y} - \hat{\alpha} - \mathbf{X} \hat{\boldsymbol{\beta}})^\top (\mathbf{y} - \hat{\alpha} - \mathbf{X} \hat{\boldsymbol{\beta}}), \end{aligned}$$

where $\bar{\mathbf{x}}$ is a vector that contains the sample means of the K predictors, \mathbf{X} is a $N \times K$ matrix of *centered* predictors, \mathbf{y} is a N -vector of outcomes, and \bar{y} is the sample mean of the outcome.

Taking a QR decomposition of the design matrix, $\mathbf{X} = \mathbf{Q}\mathbf{R}$, where $\mathbf{Q}^\top \mathbf{Q} = \mathbf{I}$ and \mathbf{R} is upper triangular, we can write the OLS solution for the coefficients as

$$\hat{\boldsymbol{\beta}} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y} = \mathbf{R}^{-1} \mathbf{Q}^\top \mathbf{y}.$$

The QR decomposition is often used for numeric stability reasons (see the familiar `lm` function in R), but, as we outline below, it is also useful for thinking about priors in a Bayesian version of the linear model.

3 Priors

The key innovation in this paper is the prior for the parameters in the QR-reparameterized model. To understand this prior, we start with the equations that characterize the maximum likelihood solutions before observing the data on \mathbf{X} and especially \mathbf{y} .

What would the prior distribution of $\boldsymbol{\theta} = \mathbf{Q}^\top \mathbf{y}$ be? We can write its k -th element as

$$\theta_k = \rho_k \sigma_y \sqrt{N-1},$$

where ρ_k is the correlation between the k th column of \mathbf{Q} and the outcome, σ_y is the marginal standard deviation of the outcome, and $1/\sqrt{N-1}$ is the standard deviation of the k column of \mathbf{Q} . Then let $\boldsymbol{\rho} = \sqrt{R^2} \mathbf{u}$, where \mathbf{u} is a unit vector that is uniformly distributed on the surface of a hypersphere. Consequently, $R^2 = \boldsymbol{\rho}^\top \boldsymbol{\rho}$ is the familiar coefficient of determination for the linear model.

An uninformative prior on R^2 would be standard uniform, which is a special case of a Beta(a, b) distribution with both shape parameters a and b equal to 1. A non-uniform prior on R^2 is somewhat analogous to ridge regression, which is popular in data mining and produces better out-of-sample predictions than least squares because it penalizes $\boldsymbol{\beta}^\top \boldsymbol{\beta}$, usually after standardizing the predictors. In our case, an informative prior on R^2 effectively penalizes $\boldsymbol{\rho}^\top \boldsymbol{\rho}$, which encourages $\boldsymbol{\beta} = \mathbf{R}^{-1} \boldsymbol{\theta}$ to be closer to the origin.

Lewandowski, Kurowicka, and Joe (2009) derives a distribution for a correlation matrix that depends on a single shape parameter $\eta > 0$, which implies the variance of one variable given the remaining K variables is Beta($\eta, \frac{K}{2}$). Thus, the R^2 is distributed Beta($\frac{K}{2}, \eta$) and any prior information about the location of R^2 , which we will denote ℓ_{R^2} , can be used to choose a value of the hyperparameter η .

Four ways of implying a value for η via the specification of ℓ_{R^2} are:

1. ℓ_{R^2} is the prior mode on the $(0, 1)$ interval.

This is only valid if $K \geq 2$ since the mode of a Beta($\frac{K}{2}, \eta$) distribution is $(\frac{K}{2} - 1) / (\frac{K}{2} + \eta - 2)$ and does not exist if $K < 2$.

2. ℓ_{R^2} is the prior mean on the $(0, 1)$ interval, where the mean of a *Betadist* $\frac{K}{2}\eta$ distribution is $(\frac{K}{2}) / (\frac{K}{2} + \eta)$.

3. ℓ_{R^2} is the prior median on the $(0, 1)$ interval.

The median of a Beta($\frac{K}{2}, \eta$) distribution is not available in closed form, but if $K > 2$ the median is approximately equal to $(\frac{K}{2} - \frac{1}{3}) / (\frac{K}{2} + \eta - \frac{2}{3})$. Regardless of whether $K > 2$, we can numerically solve for the value of η that is consistent with a given prior median.

4. ℓ_{R^2} is some (negative) prior value for $E(\log R^2) = \psi(\frac{K}{2}) - \psi(\frac{K}{2} + \eta)$, where $\psi(\cdot)$ is the Digamma function. Again, given a prior value for the left-hand side it is easy to numerically solve for the corresponding value of η .

For σ_y we set $\sigma_y = \omega s_y$ where s_y is the sample standard deviation of the outcome and $\omega > 0$ is an unknown scale parameter to be estimated. The only prior for ω that does not

contravene Bayes' theorem in this situation is the Jeffreys prior,

$$f(\omega) \propto \frac{1}{\omega},$$

which is proportional to a Jeffreys prior on the unknown σ_y ,

$$f(\sigma_y) \propto \frac{1}{\sigma_y} = \frac{1}{\omega \hat{\sigma}_y} \propto \frac{1}{\omega}.$$

This parameterization and prior makes it easy to work with any continuous outcome variable, no matter what its units of measurement are.

We need not directly specify a prior for σ_ϵ because our prior beliefs about σ_ϵ are already implied by our beliefs about ω and R^2 via the relation

$$\sigma_\epsilon = \omega s_y \sqrt{1 - R^2}.$$

Thus, the only remaining distribution to specify is a prior for $\alpha = \bar{y} - \bar{\mathbf{x}}^\top \mathbf{R}^{-1} \boldsymbol{\theta}$. As a default, an improper uniform prior is possible as the posterior will still be proper, but we have also observed that a Gaussian prior with mean zero and standard deviation σ_y/\sqrt{N} can work well in many cases.

4 Posterior

The previous sections imply a posterior distribution for ω , α , \mathbf{u} , and R^2 . After fitting the model, we can recover the parameters of interest from the primitive parameters as:

$$\begin{aligned} \sigma_y &= \omega s_y \\ \sigma_\epsilon &= \sigma_y \sqrt{1 - R^2} \\ \boldsymbol{\beta} &= \mathbf{R}^{-1} \mathbf{u} \sigma_y \sqrt{R^2 (N - 1)} \end{aligned}$$

When implementing this model, we actually utilize an improper uniform prior on $\log \omega$. Consequently, if $\log \omega = 0$, then the marginal standard deviation of the outcome *implied by the model* is the same as the sample standard deviation of the outcome. Therefore, if $\log \omega > 0$, then the marginal standard deviation of the outcome implied by the model exceeds the sample standard deviation, which is to say that the model overfits the data. If $\log \omega < 0$, then the marginal standard deviation of the outcome implied by the model is less than the sample standard deviation and so the model underfits the data (or the data-generating process is nonlinear). Given the regularizing nature of the prior on R^2 , a minor underfit would be considered ideal if the goal is to obtain good out-of-sample predictions. If the model badly underfits or overfits the data, then the model should be reconsidered.

5 Conclusion

Priors can be easy or hard for applied researchers to *specify* and easy or hard for applied researchers to *conceptualize*. Traditional shortcut priors on regression coefficients are often used because they are both easy to specify and to conceptualize. The informative prior on R^2 proposed in this paper is more difficult to conceptualize, but with the recent release of the `rstanarm` R package it is equally easy to specify.

References

- Gabry, J., and Goodrich, B. (2016). `rstanarm`: Bayesian Applied Regression Modeling via Stan. R package version 2.9.0-3. <http://cran.r-project.org/package=rstanarm>
- Guan, Y., and Stephens, M. (2011). Bayesian variable selection regression for genome-wide association studies, and other large-scale problems. *Annals of Applied Statistics*. **5**(3): 1780–1815.
- Lewandowski, D., Kurowicka D., and Joe, H. (2009). Generating random correlation matrices based on vines and extended onion method. *Journal of Multivariate Analysis*. **100**(9): 1989–2001.
- Hothorn, T., and Everitt, B. S. (2015). `HSAUR3`: A Handbook of Statistical Analyses Using R (3rd Edition). R package version 1.0-5. <http://cran.r-project.org/package=HSAUR3>
- Hothorn, T., and Everitt, B. S. (2014). *A Handbook of Statistical Analyses Using R*. Chapman & Hall/CRC, third edition.

A Example

We will utilize an example from the HSAUR3 R package, which is used in the 2014 third edition of *A Handbook of Statistical Analyses Using R*. The model in section 5.3.1 analyzes an experiment where clouds were seeded with different amounts of silver iodide to see if there was increased rainfall. This effect could vary according to covariates, which (except for time) are interacted with the treatment variable. Most people would probably be skeptical that cloud hacking could explain very much of the variation in rainfall and thus the prior mode of the R^2 would probably be fairly small.

The frequentist estimator of this model can be replicated in R by executing

```
> data("clouds", package = "HSAUR3")
> mod <- rainfall ~ seeding * (sne + cloudcover + prewetness + echomotion) + time
> ols <- lm(mod, data = clouds)
```

from which we obtain the estimated coefficients

```
> round(coef(ols), digits = 3)
      (Intercept)      seedingyes      cloudcover
      -0.346      15.683
      sne      echomotionstationary
      0.420      3.153
      prewetness      seedingyes:sne
      4.108      -3.197
      time      seedingyes:cloudcover
      -0.045      -2.557
      seedingyes:echomotionstationary
      -0.486
      -0.562
```

Note that we have *not* looked at the estimated R^2 or σ for the ordinary least squares model. We can estimate a Bayesian version of this model using the `rstanarm` package by prepending `stan_` to the `lm` call and specifying a prior mode for R^2 using the `R2` function:

```
> library("rstanarm")
> post <- stan_lm(mod, data = clouds, prior = R2(location = 0.2, what = "mode"))
```

The posterior medians for the coefficients are

```
> round(coef(post), digits = 3)
      (Intercept)      seedingyes      cloudcover
      2.303      6.753
      sne      echomotionstationary
      0.193      1.345
      prewetness      seedingyes:sne
      1.751      -1.407
      time      seedingyes:cloudcover
      -0.019
      seedingyes:echomotionstationary
```

```

-0.204                                -1.068
seedingyes:echomotionstationary
-0.264

```

The point estimates from the Bayesian model, which are represented by the posterior medians, appear quite different from the ordinary least squares estimates. However, the `log-fit_ratio` (i.e. $\log \omega$) is quite small

```

> summary(post)["log-fit_ratio", "50%"]
-0.010

```

which indicates that the model only slightly overfits the data when the prior derived above is utilized. Thus, it would be safe to conclude that the ordinary least squares estimator considerably overfits the data since there are only 24 observations with which to estimate 12 parameters and no prior information is leveraged. In general, we would expect the prior derived in this paper to be well-suited in situations like the one above, where there are many predictors relative to the number of observations.